

Extra slides about pgvector and indexing

By Tibs (they / he)

Slides and source code at
<https://github.com/aiven-labs/pgvector-find-faces-talk>

tony.ibbs@aiven.io / <https://aiven.io/tibs> / @much_of_a

Indexing

Speeds up the *use* of embeddings, but gives *approximate* results.


- IVFFlat: Inverted File with Flat Compression

Slower to search, quicker to build, smaller index, needs rebuilding

- HNSW: Hierarchical Navigable Small Worlds


Faster to search, slower to build, adapts to new data

IVFFlat: Inverted File with Flat Compression



tony.ibbs@aiven.io / <https://aiven.io/tibs> / @much_of_a

HNSW: Hierarchical Navigable Small Worlds



So which to choose?

Advice from [Vector Indexes in Postgres using pgvector: IVFFlat vs HNSW](#)

- If you care more about index size, then choose IVFFlat.
- If you care more about index build time, then select IVFFlat.
- If you care more about speed, then choose HNSW.
- If you expect vectors to be added or modified, then select HNSW.